

*Go: The Idea Behind Sync.Pool

Hang Chu

August 16, 2021

Introduction

Implementation

Benchmark

Conclusion

Introduction

- * Repeated allocation and deallocation of objects that are expensive to create and have short lifetimes is wasteful.
- * **Solution:** Create an object pool to better manage resources.
- * **Object Pool Pattern:**
 - * Initialize a pool of ready-to-use objects.
 - * Unneeded used objects are returned to the pool.
 - * If the pool is empty, a new object will be created and returned upon request.
 - * The pool must be responsible for resetting used objects' state.

- * Many ways to implement the object pooling pattern, but focus on using sync.Pool today.
- * How sync.Pool works: caching temporarily unused objects for later use.

Implementation

- * New returns a new object if pool is empty.
- * personPool has 2 methods:
 - * Get: takes an object out of the pool.
 - * Put: returns the unused object back to the pool.
- * How to measure efficiency of sync.Pool?

```
type Person struct {  
    Name string  
}  
  
// initializes pool  
var personPool = sync.Pool{  
    // returns a new object if pool is empty  
    // when called by personPool.Put()  
    New: func() interface{} { return new(Person) },  
}  
  
func main() {  
    // gets a new instance  
    newPerson := personPool.Get().(*Person)  
  
    // defers release function so the instance  
    // can be used again  
    defer personPool.Put(newPerson)  
  
    // uses the instance  
    newPerson.Name = "Jack"  
}
```

Figure 1: A simple sync.Pool example

Benchmark

Benchmarking two functions

- ✱ (De)allocation objects with and without a pool.
- ✱ Outer loop runs b.N times, with 10,000 allocations/operation (inner loop).

```
func BenchmarkWithoutPool(b *testing.B) {  
    var p *Person  
    b.ReportAllocs()  
    b.ResetTimer()  
    for i := 0; i < b.N; i++ {  
        for j := 0; j < 10000; j++ {  
            p = new(Person)  
            p.Age = 23  
        }  
    }  
}  
  
func BenchmarkWithPool(b *testing.B) {  
    var p *Person  
    b.ReportAllocs()  
    b.ResetTimer()  
    for i := 0; i < b.N; i++ {  
        for j := 0; j < 10000; j++ {  
            p = personPool.Get().(*Person)  
            p.Age = 23  
            personPool.Put(p)  
        }  
    }  
}
```

Figure 2: sync.Pool benchmark

Benchmarking two functions - Results

```
goos: linux
goarch: amd64
BenchmarkWithoutPool-4          4904          236638 ns/op      80000 B/op      10000 allocs/op
BenchmarkWithPool-4            5253          249298 ns/op         0 B/op           0 allocs/op
PASS
ok      _/home/jodi/tfs-03/lec-06/presentation 4.294s
```

Figure 3: Benchmark results

- ✱ BenchmarkWithoutPool took 80,000B to allocate 10,000 objects per operation, while BenchmarkWithPool didn't take any.
- ✱ Not the best example to show the disadvantages of sync.Pool due to negligible difference in running time per operation between two functions.

No such thing as a free lunch

- * Test performance of sync.Pool in a parallel setting.
- * The sync.Pool version took **16X** more time per operation to run than the allocation one.

```
func BenchmarkPool(b *testing.B) {  
    var p sync.Pool  
    b.RunParallel(func(pb *testing.PB) {  
        for pb.Next() {  
            p.Put(x 1)  
            p.Get()  
        }  
    })  
}  
  
func BenchmarkAllocation(b *testing.B) {  
    b.RunParallel(func(pb *testing.PB) {  
        for pb.Next() {  
            i := 0  
            i = i  
        }  
    })  
}
```

Figure 4: Benchmarking sync.Pool and simple Allocation

BenchmarkPool-4	100000000	10.9 ns/op
BenchmarkAllocation-4	1000000000	0.674 ns/op

Figure 5: Benchmarking results

How sync.Pool works under the hood

- * sync.Pool has 2 containers: local pool and victim cache.
- * package sync registers init function to the runtime to clean the pool, which will be triggered by the garbage collector
- * When GC is triggered, victim cache is cleared and objects from local pool will be moved there.
- * Both new and returned objects are put in local pool.

```
func init() {  
    runtime_registerPoolCleanup(poolCleanup)  
}
```

Figure 6: sync package's init function

```
func poolCleanup() {  
    // Drop victim caches from all pools.  
    for _, p := range oldPools {  
        p.victim = nil  
        p.victimSize = 0  
    }  
  
    // Move primary cache to victim cache.  
    for _, p := range allPools {  
        p.victim = p.local  
        p.victimSize = p.localSize  
        p.local = nil  
        p.localSize = 0  
    }  
  
    oldPools, allPools = allPools, nil  
}
```

Figure 7: Pool cleanup function

Conclusion

- * Need to create expensive objects a lot of times? Use object pooling, more specifically `sync.Pool`.
- * Object pooling is slower than simple (de)allocation under certain conditions.